
Self-Hosted Agents

Contents

Intro.....	2
The Simplest Real-World Business Setup	3
One Dedicated PC or VM	3
Connect Only a Few Things	4
Use Case #1: Internal Knowledge Assistant	4
Business Benefit	5
Use Case #2: Shared Inbox Assistant	5
Business Benefit	5
Use Case #3: SOP Execution Assistant	5
Use Case #4: IT Help Desk Assistant	5
What I Would NOT Do Initially	6
My Recommended Architecture.....	6
If You Are a Microsoft Shop	7
Maturity Model.....	7
Level 1: Proof of Concept (Free)	7
Level 2: Departmental Solution	8
Level 3: Serious Internal AI Platform	9
What Most Production Shops Actually Do.....	9
80% Search.....	9
15% Automation	10
5% Agentic Stuff	10
The Natural "Next Step Up"	10
Front End.....	10
Model	10
Knowledge.....	10
Automation	10
Identity.....	10
Documents	10

Storage	11
The Executive Conversation.....	11
Phase 1	11
Phase 2	11
Phase 3	11

Intro

One of the quickest free fully self-hosted options right now is SafeAgent (open-source, Docker-based). It runs entirely on your machine, includes a local model option (no external API keys required to start), and gives you a chat interface with real agent capabilities (tools for web, email, files, workflows, memory, etc.).

[Get Started | Docker](#)

[SafeAgent — Your AI, Your Machine, Your Data](#)

Fastest start (≈30–60 seconds if Docker is already installed)

```
curl -sL https://www.safeagent.dev/docker-compose.yml -o docker-compose.yml && docker compose up
```

Then open <http://localhost:3000>, create an account, and start chatting. It ships with a local 7B model via Ollama for zero ongoing cost and full offline use. You can optionally connect external models or services later.

Other very quick free/self-hosted routes

GoGogot (ultra-lightweight, 15 MB binary, 10 MB RAM idle)

One Docker command + Telegram bot. Excellent for a personal always-on agent with shell, files, web, memory, and scheduling.

Uses free/cheap models (e.g. DeepSeek via OpenRouter at ~\$0.02/session) or your own providers. Ideal if you want something that lives on a cheap VPS and talks to you on Telegram.

Ollama + Open WebUI or AnythingLLM (fully local, classic stack)

Install Ollama (`curl -fsSL https://ollama.com/install.sh | sh` or the native installer).

`ollama pull qwen2.5` (or another tool-calling model like `llama3.1` / `mistral`).

Run Open WebUI:

```
docker run -d -p 3000:8080 --add-host=host.docker.internal:host-gateway -v open-webui:/app/backend/data --name open-webui --restart always ghcr.io/open-webui/open-webui:main
```

Or AnythingLLM (also Docker, strong RAG + agent mode with @agent).

Both give you a ChatGPT-like UI with local models and agent/tool features. Completely free once the model is downloaded.

n8n (visual workflow agent builder)

```
docker run -it --rm --name n8n -p 5678:5678 -v ~/.n8n:/home/node/.n8n n8nio/n8n
```

Connect it to a local Ollama instance and build agent-style automations with a drag-and-drop UI. Great if you prefer no-code workflows.

Tips for staying fully free

Prefer local models via Ollama (or the ones bundled in SafeAgent/LocalAI) so you pay \$0 for inference.

Tool-calling models that work well locally in 2026 include Qwen2.5/Qwen3 variants, Llama 3.1/3.3, and similar mid-size models (8B–30B depending on your RAM/GPU).

Hardware: 8–16 GB RAM is comfortable for smaller models; more helps with larger ones or concurrent use.

All of the above are open-source and self-hosted—no vendor lock-in or forced cloud dependency.

SafeAgent is currently one of the lowest-friction “one command and you’re chatting with a real agent” options that stays free and private by default. If you tell me your OS, hardware constraints, or preferred interface (web UI, Telegram, terminal, etc.), I can refine the exact steps.

If I were advising a **small to mid-sized business** that wants a **simple, low-cost, private AI agent**, I would avoid trying to make SafeAgent a company-wide Copilot replacement on day one. Instead, I'd use it as a **local operations assistant** for a few targeted business workflows.

SafeAgent is designed as a self-hosted AI agent that runs locally, includes a chat interface, supports local models via Ollama, can connect to email, Slack, GitHub and document repositories, and keeps processing on infrastructure you control. [\[safeagent.dev\]](https://safeagent.dev), [\[safeagent.dev\]](https://safeagent.dev)

The Simplest Real-World Business Setup

One Dedicated PC or VM

Hardware

- Windows 11 PC or Ubuntu VM
- 16 GB RAM (minimum I'd recommend)

- Docker Desktop
- SafeAgent
- Ollama with a small model (Qwen or Llama)

This becomes your "Business AI Assistant Server."

Connect Only a Few Things

Connect:

☒ Shared mailbox (support@, sales@, info@)

☒ Shared file location

- SharePoint sync folder
- Network drive
- Policy documents
- SOPs

☒ Teams/Slack notifications (optional)

Avoid connecting:

- ERP
- Finance systems
- Production databases

until you've proven value.

Use Case #1: Internal Knowledge Assistant

This is probably the easiest win.

Upload:

- Policies
- Procedures
- Employee handbook
- Product documentation
- Training guides

Employees ask:

"How do I onboard a new employee?"

"What's our PTO policy?"

"Where is the expense reimbursement form?"

SafeAgent's RAG/knowledge base capability can answer from your uploaded documents rather than general internet knowledge. [\[safeagent.dev\]](#), [\[safeagent.dev\]](#)

Business Benefit

- Less interruption of managers
 - Faster answers
 - Data stays local
-

Use Case #2: Shared Inbox Assistant

Connect a departmental mailbox.

Examples:

"Summarize today's customer emails."

"Show me all emails mentioning pricing."

"Draft responses for unanswered emails."

SafeAgent supports email reading and sending through standard mail services according to its documentation. [\[safeagent.dev\]](#), [\[safeagent.dev\]](#)

Business Benefit

- Customer service triage
 - Sales lead sorting
 - Reduced admin work
-

Use Case #3: SOP Execution Assistant

Employees can ask:

"Generate a customer onboarding checklist."

"Create a project kickoff plan."

"Summarize this PDF."

"Extract action items from this document."

This is often where local agents provide immediate value without requiring deep integrations.

Use Case #4: IT Help Desk Assistant

Feed it:

- Network documentation
- Server procedures

- User guides
- Internal KB articles

Then staff can ask:

"How do I reset VPN access?"

"What's the printer setup procedure?"

"What are the steps for laptop provisioning?"


What I Would NOT Do Initially

I would avoid:

- ✗ Allowing autonomous email sending
- ✗ Giving admin rights
- ✗ Connecting directly to production databases
- ✗ Letting the agent execute scripts unattended

Even though SafeAgent includes workflow and action capabilities, keeping it mostly "read and summarize" for the first phase significantly reduces risk. The platform itself emphasizes approvals and risk-scoring for actions. [\[safeagent.dev\]](#), [\[safeagent.dev\]](#)

My Recommended Architecture



No cloud AI.

No per-user licensing.

No API charges.

No sensitive data leaving your environment when using local models. SafeAgent states it can run with local AI models and without external API keys. [\[safeagent.dev\]](#), [\[safeagent.dev\]](#)

If You Are a Microsoft Shop

Honestly, the sweet spot for someone like you who already works heavily with M365 would be:

SafeAgent + SharePoint Library + Shared Mailbox + Teams notifications

That gives you:

- Internal knowledge search
- Email triage
- Document summarization
- Workflow assistance
- Local/private AI

without trying to replace Microsoft 365 or Power Automate.

Think of it as a **private junior operations analyst** sitting on one internal server rather than an enterprise AI platform.

For a first pilot, I'd start with **5-10 users, one shared mailbox, and a SharePoint-synced document library**. That's usually enough to demonstrate value in a week.

Absolutely. If the goal is "**show me the art of the possible without spending money**", SafeAgent is actually a pretty good demo platform because people immediately understand the value of:

- Chat with company documents
- Email triage
- Task execution
- Local/private AI
- No subscription

The challenge is that the moment leadership starts saying "*we want 50 users on this*" you've crossed from a **lab/demo architecture** into a **production architecture**.

Maturity Model

Level 1: Proof of Concept (Free)

1	SafeAgent
2	
3	Ollama
4	
5	Single PC / VM

Users:

- 1 to 5

Purpose:

- Demonstrate concepts
- Internal KB
- Shared mailbox

Cost:

- Essentially free


Risk:

- Low

This is where most AI pilots should start.

Level 2: Departmental Solution

At this point you stop thinking about "my AI" and start thinking about a service.



Changes:

- Dedicated VM or server
- Nightly backups
- Entra ID authentication
- Shared knowledge repositories
- Teams integration
- Monitoring

Users:

- 10 to 50

Cost:

- A few hundred dollars of hardware


This is usually the point where management says:

"This is actually useful."

Level 3: Serious Internal AI Platform

This is where most organizations end up.

Instead of SafeAgent being the centerpiece:



Typical stack:

- Open WebUI
- Ollama or vLLM
- Qdrant/Chroma/Pinecone
- n8n
- PostgreSQL
- Entra ID SSO

SafeAgent may still exist, but now it's one component in a platform.

What Most Production Shops Actually Do

This is the interesting reality.

Most companies discover they don't need:

- autonomous agents
- complex planning systems
- AI making decisions

What they need is:

80% Search

"Find me information."

15% Automation

"Create a ticket."

"Draft a response."

"Summarize this meeting."

5% Agentic Stuff

"Go perform multiple actions for me."

That means a surprisingly effective architecture is:

1	SharePoint
2	Teams
3	Exchange
4	
5	RAG Layer
6	
7	Local LLM
8	
9	Chat Interface

Very little autonomy.

Very high business value.

The Natural "Next Step Up"

If I had to recommend one realistic production open-source stack today for an M365 organization, it would probably be:

Front End

Open WebUI

Model

Ollama or vLLM

Knowledge

Qdrant

Automation

n8n

Identity


Entra ID

Documents

SharePoint Synced Libraries

Storage

PostgreSQL



Why?

Because every piece can later be swapped out.

No lock-in.

The Executive Conversation

The conversation I would have with leadership is:

Phase 1

"We can stand up an AI agent internally this week for essentially no software cost."

Phase 2

"We'll connect it to a non-sensitive document set and a shared mailbox."

Phase 3

"If employees actually use it, we can evolve it into a supported internal AI platform."

That approach changes the discussion from:

"Should we buy a \$50,000 AI solution?"

to

"Should we invest a few days proving whether this saves any time at all?"

That's usually a much easier sell.

And honestly, Tim, for a Microsoft-heavy shop, the most compelling demo is not "Look, I built an agent."

It's:

"Here's 10,000 SharePoint documents, a support mailbox, and a local AI that answers questions in 5 seconds without leaving our network."

Business leaders immediately understand that use case.