
Self-Hosted Agents

One of the quickest free fully self-hosted options right now is SafeAgent (open-source, Docker-based). It runs entirely on your machine, includes a local model option (no external API keys required to start), and gives you a chat interface with real agent capabilities (tools for web, email, files, workflows, memory, etc.).

[Get Started | Docker](#)

[SafeAgent — Your AI, Your Machine, Your Data](#)

Fastest start (≈30–60 seconds if Docker is already installed)

```
curl -sL https://www.safeagent.dev/docker-compose.yml -o docker-compose.yml && docker compose up
```

Then open <http://localhost:3000>, create an account, and start chatting. It ships with a local 7B model via Ollama for zero ongoing cost and full offline use. You can optionally connect external models or services later.

Other very quick free/self-hosted routes

GoGogot (ultra-lightweight, 15 MB binary, 10 MB RAM idle)

One Docker command + Telegram bot. Excellent for a personal always-on agent with shell, files, web, memory, and scheduling.

Uses free/cheap models (e.g. DeepSeek via OpenRouter at ~\$0.02/session) or your own providers. Ideal if you want something that lives on a cheap VPS and talks to you on Telegram.

Ollama + Open WebUI or AnythingLLM (fully local, classic stack)

Install Ollama (`curl -fsSL https://ollama.com/install.sh | sh` or the native installer).

`ollama pull qwen2.5` (or another tool-calling model like llama3.1 / mistral).

Run Open WebUI:

```
docker run -d -p 3000:8080 --add-host=host.docker.internal:host-gateway -v open-webui:/app/backend/data --name open-webui --restart always ghcr.io/open-webui/open-webui:main
```

Or AnythingLLM (also Docker, strong RAG + agent mode with @agent).

Both give you a ChatGPT-like UI with local models and agent/tool features. Completely free once the model is downloaded.

n8n (visual workflow agent builder)

```
docker run -it --rm --name n8n -p 5678:5678 -v ~/.n8n:/home/node/.n8n n8nio/n8n
```

Connect it to a local Ollama instance and build agent-style automations with a drag-and-drop UI. Great if you prefer no-code workflows.

Tips for staying fully free

Prefer local models via Ollama (or the ones bundled in SafeAgent/LocalAI) so you pay \$0 for inference.

Tool-calling models that work well locally in 2026 include Qwen2.5/Qwen3 variants, Llama 3.1/3.3, and similar mid-size models (8B–30B depending on your RAM/GPU).

Hardware: 8–16 GB RAM is comfortable for smaller models; more helps with larger ones or concurrent use.

All of the above are open-source and self-hosted—no vendor lock-in or forced cloud dependency.

SafeAgent is currently one of the lowest-friction “one command and you’re chatting with a real agent” options that stays free and private by default. If you tell me your OS, hardware constraints, or preferred interface (web UI, Telegram, terminal, etc.), I can refine the exact steps.